

RoleUp demo apps User Guide Package

User guide for the matched expert and learner set


Internal · August 2026 · files: RoleUp_Expert_Demo_anonymized_v2.html and
RoleUp_Learner_Demo_anonymized.html

Contents:

RoleUp demo apps User Guide Package	1
The matched set, in one picture.....	2
The expert app: Jordan.....	2
The learner app: Riya.....	5
Q&A.....	9
On validity.....	9
On the method and the relationship	9
On the evidence.....	9
On healthcare and data.....	10
The one-sentence answers, if time is short.....	10
Walkthrough Instructions.....	10
Practical notes	10
RoleUp glossary	11

The matched set, in one picture

The two apps are one product photographed from both ends. Jordan Reyes (JR) is the expert; his app is the capture surface that has been running since June. Riya Patel (RP) is the learner; her app is the dispense surface, prototyped this week. The engine between them is the same diff run in two directions, and the demo's payoff is watching a single case travel the full loop: a contradiction leaves Riya's phone and arrives on Jordan's as an elicitation probe with a merge dossier attached.




One engine, two directions: the simulation interview, both ways

Both files open in any browser, run fully offline and hold no external dependencies. Reloading the page resets all state, which is the fast recovery if a walkthrough goes sideways. Tap targets behave like a normal mobile app.

The expert app: Jordan


Persona: Jordan Reyes, director of software engineering at Cedarview Health Network, day 15 of capture, 24 claims across 5 responsibilities. Everything on this side demonstrates capture with practitioner sovereignty; the learner loop is the v2 addition that closes the circle.



Home

Jordan's daily surface. "Pin it" is the primary action: a seconds-long capture for the moment something just happened, before the reasoning evaporates. "Check-in" is the short morning session (shown done at 7:48 AM) and "Percolate" holds 4 pins ready to walk through.


Two nudge cards follow: 3 claims ready for release review, and the learner loop card showing 1 elicitation probe and 1 merge dossier from Riya Patel. The bottom list, "What Clara is learning", opens each responsibility with its claims, each tagged by knowledge category and a freshness chip (fresh, aging, stale). That chip is the primitive form of decay, worth pointing at when the drift sensor comes up.



Percolate: the cognitive session

Clara opens with a sentiment step, then walks yesterday's pins one at a time, probing the reasoning behind each before structuring it into claims. Responses are tappable.


This is the elicitation engine on screen. Notice the revalidation note in coral: Clara is tracking a survival-wisdom claim that has not been confirmed in 11 days and offers to re-test it. Her questions always aim at what she does not yet understand, which is the capture-direction diff in action.



Review and release

Sovereignty made visible, and the screen to slow down on with any audience. Claims sit in two groups, ready (3) and released (21). Each card shows its category, freshness and full text; "Release to org" is Jordan's button and nothing moves without it.


The principle box states the standing rule: RoleUp surfaces, Jordan decides. Released claims move to the released group with a confirming toast.



Learner loop: the probe

Opened from the home nudge. The header carries the thesis: your graph is teaching, and teaching is testing the graph.

The elicitation probe reports that Riya hit a vendor-escalation case Jordan's model does not cover, because the July 2026 provincial regulation postdates his May claim. Her confusion arrives as targeted elicitation fuel, and "Walk it through with Clara" queues the case for his next percolate.



Learner loop: the merge dossier


Riya's candidate claim from the reviewer's seat, with everything Clara prepared: lineage back to Jordan's two-cycle rule, 3 triggering cases, a backtest that held in 14 of 16 historical cases and a passed synthetic cast replay. Clara did the investigation; the judgment is Jordan's.

The five merge verbs appear as chips with supersedes proposed, scoped to regulated contexts, his rule staying canon everywhere else. "Merge" and "Decline & capture why" both resolve with toasts; the decline path shows rejection becoming negative knowledge rather than a dead end. The principle to read aloud: the reviewer is the role, never the person.

Also worth showing if time allows: Tools holds the week 4 outputs, an onboarding brief generated from 24 patterns that refreshes as the graph grows, which is the first dividend of capture and the natural bridge into the learner story.

The learner app: Riya


Persona: Riya Patel, software engineer, week 6, who has cloned the director of software engineering role graph that Jordan authored. This side demonstrates the dispense direction: the graph as answer key, the learner as the unknown. It shares Jordan's stylesheet so the pair reads as one product.



Home

The hero is the diff itself: 38% of the role's required judgment demonstrated, with trend and streak. This number is diff closure, the metric the whole utility argument rests on. A coral nudge tracks her open contradiction with Jordan.


Below, "Your diff, by region" lists five regions with closure bars and posture chips. Four read Mentor. Incident triage, at 82%, reads "Capturing from you" in gold: Clara's posture flipped there because Riya demonstrated the region. That one chip is the whole posture-per-region argument; there is no graduation day.



Rehearse: the scenario

Scenario 1, legacy data plumbing: Monday 8:40 AM, warehouse tables stale, vendor dashboard green. The card notes it was built from 3 claims with Jordan's captured cues as the answer key.


Three choices, and the wrong ones are worth picking on purpose: "open a vendor ticket" or "re-run the sync" trigger route one, reinforcement, where Clara names the discriminating cue and re-presents the choice. It proves the system teaches rather than just grades.



The cue probe

Clara says right call, then asks the question that separates knowing the answer from holding the judgment: what told you?


Answering with the cue (green dashboard plus stale tables means the gap is downstream) moves the edge to demonstrated. Answering "I remembered someone mentioned an ETL" moves it only to practiced, and Clara says she will come back to it. Either way a ledger update card appears with full provenance: Jordan Reyes, June 23, 2026, his words quoted, lineage attached.



Contradiction routing

Scenario 2 is her thinnest region and the demo's centerpiece. The canon on the card is Jordan's May two-cycle rule; choosing "escalate now, flagged as data integrity" cites the July regulation that postdates it.


Clara shows the routing card with all three routes and route three lit: the learner may be right because the world moved. The status block confirms the case went to Jordan as an elicitation probe, a dossier is being prepared and her ledger marks the edge contradicted. Clara's closing line carries the drift-sensor thesis: either way this resolves, the graph comes out sharper.



The session diff

The session ends with what changed, not a transcript: state movements, the cue recognized, the contradiction routed and the next target named, with the note that only aggregate closure is shared, with consent.


"Wrap up" returns home, where the diff ticks from 38% to 41% with a toast. That moving number is transfer being measured, and it is the moment to hold for a beat.



The ledger

The six-state legend runs across the top: 46 unencountered, 31 encountered, 18 questioned, 22 practiced, 23 demonstrated, 1 contradicted. Region cards below show each region's distribution.

The framing line matters: the graph stays single-source; this overlay is hers. Her lead sees aggregate closure, never her stumbles, unless she shares.



The merge gate

Scrolling down reaches her candidate claim in full: the diverging text, lineage to Jordan's rule, the Clara-prepared dossier and the five verbs with supersedes proposed.

Two lines to point at: if declined, Clara captures the reasoning and the claim stays in her graph; and under "Your emerging graph", the attribution principle, the claim travels with the person, forever. Ask and Profile round out the app: sample graph queries with provenance, the gold-versus-Jordan learning source selector, and the standing rule that she is an author in training, not a subject.

Q&A

Talking points for demo.

On validity

1. **How do you validate that machine-generated scenarios preserve the simulation interview's validity?** We don't claim they do yet; that's precisely the question we want examined. The scenarios are generated from captured cues and expectancies rather than invented, so they inherit their content validity from the elicitation. Whether the generation preserves the properties that make the simulation interview work is, as far as we can tell, a question never asked in the literature, because no instrument could pose it before. That's the research object the Design stream exists for.
2. **Your answer key is a snapshot. What happens when the graph is wrong and a learner gets trained on it anyway?** Three things, by design. Contradiction routing means a learner whose reasoning beats the canon triggers review rather than correction, so the system has a path for the graph being wrong. Provenance means every claim carries its author and capture date, so staleness is inspectable rather than hidden. And decay tracking means unvalidated claims surface for re-confirmation. The failure mode you're describing, confidently teaching stale judgment, is the default failure of every knowledge repository ever built. Ours is the first architecture we know of that treats it as a first-class event with a routing path.
3. **How is reasoning-trace assessment scored without a human rater?** Against structure, not against prose. The expert's captured model gives us the discriminating cue, the expectancy and the chosen strategy for a decision point. The learner's response either names the cue or doesn't, matches the expectancy or doesn't. We're comparing a trace to a formal reference, which is a narrower and more tractable task than open-ended judgment of quality. Where it gets genuinely hard, partial matches, novel-but-valid reasoning, is exactly where the contradiction path routes to a human. The machine scores the clear cases; the interesting cases escalate.
4. **Isn't the synthetic cast circular? You're validating an LLM system with an LLM.** For elicitation coverage, yes, which is why it's benchmarked against the NASA tacit knowledge capture transcripts: twenty public-domain human sessions we can compare against for what a real elicitation surfaces. The cast is a regression harness and a stress test, not a proof. Human-parity claims wait for human data, and the pilots are where that comes from.

On the method and the relationship

5. **How is the mentor architecture different from ShadowBox?** It's the closest living relative, and we'd rather say so than have you conclude it. Two differences. ShadowBox scenarios are authored by experts and researchers; ours are generated from a formal graph captured from the specific expert whose judgment is being transferred, so the scenario set updates when the graph does. And ShadowBox compares a trainee's choices to an expert panel's; we compare a reasoning trace to a per-claim reference with provenance, which yields a per-node ledger rather than a session score. Whether those differences preserve what makes scenario-based training work is a question we'd rather investigate with the people who built it than around them.
6. **You've encoded our tradition's methods in a commercial product and filed patents. What's your position on that?** The patents cover our system: the routing, the representation, the diff mechanics, the sovereignty architecture. They don't and can't claim ACTA, CDM or the simulation interview, which are published science we cite by name in our materials, including the ones investors read. The category paper credits the lineage explicitly. What we've built is an instrument that runs the method; the method remains yours, published, attributed. If any wording anywhere reads otherwise to you, we want to know and we'll fix it.
7. **What do you actually want from ADS?** Three things, in the shape we've discussed. Verify: arm's-length recurring review of whether Clara's elicitation meets the published method's standard, and now whether the inversion preserves its validity. Design: co-development on the questions neither of us can answer alone, with a publication path. Scale: Clara configurations for ADS delivery where that serves your work. And an independence firewall stated plainly: your review conclusions are yours, whatever they say.

On the evidence

8. **Your capture engagement data shows abandoned sessions. Why should transfer engagement be different?** It might not be, and we've published that risk to our own investors. Two structural differences give us hope rather than confidence: a learner has a sharper daily incentive than an expert (they need the

judgment now, the expert already has it), and rehearsal produces a visible personal metric, the closing diff, where capture's payoff is deferred and diffuse. Grudin's law says collaborative systems fail when cost and benefit are asymmetric; dispense is the side of our system where cost and benefit land on the same person. The twelve-practitioner engagement and the pilots will tell us.

9. **What would make you conclude the transfer thesis is wrong?** If rehearsal against the graph doesn't close the diff faster than shadowing and documentation do, measured in the pilots, the thesis fails and we're a capture and consultation company. We've written that sentence into our investor materials. Diff closure is a designed instrument, not a result; September's validation program and the pilot Phase 2s are the dated tests.

On healthcare and data

10. **How does patient-adjacent data stay out of the graph?** Capture is about the practitioner's reasoning, not case content, and the extraction layer structures decisions, cues and strategies rather than records. Deployment is customer-hosted where required, the practitioner gates every release, and validating cases in the authorship path stay classified with the organization; only the decision rule and an attestation travel. Healthcare is our beachhead precisely because we designed for its constraints rather than retrofitting them.

The one-sentence answers, if time is short

11. Why now: conversational AI made a thirty-year-old method affordable; someone will industrialize it, and it should be done to the method's standard.
12. What's running: capture, in production.
13. What's prototyped: the full learner loop, on two phones in this room.
14. What's next: the production ledger, after capture stabilizes.
15. What we're asking: hold the instrument to the standard of the people who built the method.

Walkthrough Instructions

Run Riya first, produce the contradiction deliberately, then switch phones so the same case arrives on Jordan's side as the payoff. The full loop takes about four minutes at demo pace.

One. Riya's home: linger on the diff number and the posture chips, especially incident triage in gold.

Two. Rehearse: pick a wrong answer first to show reinforcement, then the right action, then the cue answer so the edge moves to demonstrated with Jordan's provenance attached.

Three. Scenario 2: choose the divergent escalation so the contradiction routes. Read route three aloud.

Four. Wrap up: the session diff, the toast and the number moving to 41%.

Five. Riya's ledger: the six states, then scroll to the merge gate card and the five verbs.

Six. Switch to Jordan's phone: home nudge, then the learner loop. The probe first, then the dossier. Merge it, or decline it and show that rejection is captured too.

Practical notes

- Both apps are single self-contained files: no network, no install, they run from a downloads folder.
- Reloading resets all state, so a botched take costs ten seconds.
- The rehearsal branches; every path is designed, so there is no wrong tap, but the walkthrough above is the path that shows the most.
- Have both phones pre-loaded and run the loop once end to end before the meeting.
- Claims discipline when presenting:
 - the capture app reflects what runs in production today;
 - the learner surface is the design made operable, prototyped this week and sequenced next.
 - Present it as exactly that.

RoleUp glossary

Internal, canonical. One definition per term; every document, demo and conversation inherits these. Register notes mark where a term is market-facing (M), internal (I) or patent language (P).

1. **Judgment capture** (M). The category. Capturing how an organization's most experienced people decide, through structured conversation in the flow of work, into a permanent, attributed asset. Our coinage.
2. **Judgment continuity** (M). The buyer's problem the category solves: succession, onboarding and the retirement wave. Used in market-facing materials where "capture" names the act and "continuity" names the outcome.
3. **Organizational logic graph, OLG** (M). The formal representation of captured judgment: decision points, cues, expectancies and strategies, plus the workflows, relationships, exceptions and systems judgment runs on, with provenance on every claim. The same structure appears as the institutional logic graph, ILG, in patent language (P).
4. **Clara** (M). The elicitation agent. Conducts sessions, extracts structure per turn, generates scenarios, prepares dossiers. Clara surfaces; humans decide.
5. **Claim** (I). The atomic unit of the graph: one piece of judgment with its verbatim quote, session citation, author, category and freshness.
6. **Pin** (M). A seconds-long capture of a moment worth returning to, made in the flow of work. Raw material for percolation.
7. **Percolate** (M). The session where Clara walks pinned moments through with the practitioner, probing the reasoning and structuring it into claims.
8. **Three runs** (I). The methodology's passes per role: designed intent (run 1, document ingestion), enacted practice (run 2) and the rationale for the gap between them (run 3).
9. **Elicitation technique library, ETL** (I). Ten research traditions encoded as a machine-readable, provenance-tracked graph; the session planner routes between probe families per the coverage state.
10. **Capture context** (I). The defensibility point: selection-grade evidence of how someone decides, gathered over weeks while nothing is at stake, versus competitors harvesting performance at maximum stakes.
11. **Practitioner sovereignty** (M). The architectural rule that practitioners are authors, not subjects: they gate every release, remain named, and the surveillance version of the product is structurally absent.
12. **Release gate** (I). The mechanism of sovereignty: nothing reaches the organization's graph until the practitioner reviews and releases it.
13. **Diff engine** (I). The platform's one mechanism: compare a reference to an unknown, work on the gap. Capture points it at the expert (reference: their mind; unknown: Clara's model). Dispense inverts it (reference: the graph; unknown: the learner).
14. **Dispense** (I). The transfer direction: scenarios generated from the graph, the expert's captured cues and expectancies as answer key, the learner's reasoning trace assessed against them.
15. **Diff closure** (M). The transfer metric: the computable gap between a role's required judgment and a learner's demonstrated judgment, closing. The number the utility argument rests on.
16. **Learner ledger** (I). A per-learner overlay on the OLG, never a copy, sovereign to the learner. Records one of six states per node or edge in scope.

17. **Six states** (I). Unencountered, encountered, questioned, practiced, demonstrated, contradicted. The ledger's vocabulary for a learner's relationship to each piece of the graph.
18. **Posture** (I). Which direction Clara points with a given person in a given graph region: mentor (teaching) or elicitation (capturing). Scoped per region, never per person; a demonstrated region flips Clara from mentor to capture. There is no graduation day.
19. **Contradiction routing** (I). The three-way fork when a learner's reasoning disagrees with the canon: learner wrong (reinforcement), graph incomplete (elicitation probe to the expert), learner right (authorship trigger).
20. **Drift sensor** (I). What routing makes the graph: divergence surfaces staleness through use, with evidence attached, instead of letting it accrue silently. Every cohort that learns from the graph leaves it sharper.
21. **Dossier** (I). Clara's prepared package for review of a divergent claim: lineage, triggering cases, backtest against historical cases, synthetic replay. The reviewer judges a dossier, not a raw claim.
22. **Merge gate** (I, horizon). The review mechanism for candidate claims. The reviewer is the role, never the person. Horizon material: verbal only outside NDA contexts.
23. **Five merge verbs** (I, horizon). Supersedes, specializes, generalizes, extends, contradicts. Merge is a vocabulary, never a binary; priority is conditional on context, never scalar over authors.
24. **Negative knowledge** (I). A declined claim captured with its reasoning: why we still do not do it that way. Rejection enriches the graph.
25. **Lineage** (I). The attribution chain a claim carries: who authored it, what it diverged from, when it merged. The practitioner's citation. Inalienable; travels with the person.
26. **Abstraction boundary** (I, horizon). The sovereignty wall at authorship: the claim (portable, the author's) versus the validating cases (classified, the organization's), bridged only by attestation, a signed statement that validation occurred without any case leaving.
27. **Synthetic cast** (I). Sixteen synthetic practitioners across five verticals, benchmarked against the NASA tacit knowledge capture transcripts. Standing infrastructure: assessment QA, regression testing, graph-health monitoring, merge validation.
28. **Gold version** (I). The consensus projection of a role's graph, versus an individual expert's own version. Learners take gold for essentials and an expert's version for depth; the quirks are where the judgment lives.
29. **Transfer velocity** (M). Diff closure gained per rehearsal hour.
30. **Succession readiness** (M). Whether anyone else has demonstrated the judgment a role depends on, as a computed number rather than an attestation. ISO 30414 names the category; we supply the arithmetic.
31. **Bottleneck decompression** (M). Interruptions deflected from the expert because the answer was reachable without them.
32. **PLG** (I). The practitioner's own graph, the default architecture and the sovereignty substrate; organizational data flows off it only on the practitioner's terms.
33. **Awen** (I). The solo-practitioner surface running on the same graph and elicitation engine, built separately.

Claims discipline, standing: capture, ledger and measurable diff closure are near-term and may appear in writing. Merge gate, authorship economics and the marketplace return path are horizon: verbal only, except under NDA with design partners.